

Deep Learning — Assignment 6

Assignment for week 6 of the 2022 Deep Learning course (NWI-IMC070) of the Radboud University.

Names: Luka Mucko, Luca Poli

Group: 46

Instructions:

- Fill in your names and the name of your group.
- Answer the questions and complete the code where necessary.
- Keep your answers brief, one or two sentences is usually enough.
- Re-run the whole notebook before you submit your work.
- Save the notebook as a PDF and submit that in Brightspace together with the `.ipynb` notebook file.
- The easiest way to make a PDF of your notebook is via File > Print Preview and then use your browser's print option to print to PDF.

Objectives

In this assignment you will

1. Build a graph neural network, using pytorch geometric
2. Compare a GNN with other network architectures
3. Compare different GNN layers and aggregation functions

Required software

As before you will need these libraries:

- `torch`, `torch-sparse`, `torch-scatter`, and `torch-geometric` for PyTorch,
- `d2l`, the library that comes with [Dive into deep learning](#) book.

The recommended way to install these libraries is described in the [torch-geometric installation instructions](#).

```
In [22]: # Replace ${TORCH} and ${CUDA} with your torch and cuda versions.  
# Or remove the -f argument to compile from source  
#
```

```
#!pip install torch-scatter torch-sparse -f https://data.pyg.org/whl/torch-${TORCH}
#!pip install torch-geometric
```

```
In [23]: %config InlineBackend.figure_formats = ['png']
%matplotlib inline

from d2l import torch as d2l
import itertools
import numpy as np
import matplotlib.pyplot as plt
import torch
import torch_geometric
from torch import nn
from torch.nn import Linear, Dropout
from torch.nn import functional as F
from torch_geometric.datasets import Planetoid
from torch_geometric.transforms import NormalizeFeatures
from torch_geometric.nn import GCNConv, SAGEConv, GraphConv
```

6.1 A node classification dataset (1 point)

In this assignment we will be working on a node classification problem using the Citeseer dataset. This is a graph dataset that contains bag-of-words representation of documents and citation links between the documents. So there is an edge between document i and document j if one cites the other. This is an undirected edge.

```
In [24]: dataset = Planetoid(root='data', name='Citeseer', transform=NormalizeFeatures())
```

(a) How many graphs are there in this dataset? How large are they (in terms of nodes and edges)? **(1 point)**

```
In [25]: # TODO: your answer here
print(f'Number of graphs: {len(dataset)}')
print(f"Number of nodes: {dataset[0].num_nodes}")
print(f"Number of edges: {dataset[0].num_edges}")
```

```
Number of graphs: 1
Number of nodes: 3327
Number of edges: 9104
```

In fact, we will continue the rest of this notebook using the first graph from the dataset.

```
In [26]: data = dataset[0] # Get the first graph object.
```

We will be use a subset of the nodes for training, and another subset for testing. These subsets are indicated by `data.train_mask` and `data.test_mask` respectively.

6.2 MLP for node classification (6 points)

In theory, we should be able to classify documents based only on their content, that is, using the bag-of-words features, without taking the graph structure into account.

We can verify that by constructing a simple node-wise multilayer perceptron with a single hidden layer. This network does not use the edge information at all.

(a) Complete the code below.

(2 points)

The network should have 2 linear layers. The hidden layer should have size `hidden_channels`, use ReLU activations, and use dropout with a dropout rate of 0.1. Don't use an activation function after the final layer.

Hint: avoid using `Sequential`, it will make the assignment harder later on.

```
In [27]: class MLP(torch.nn.Module):
    def __init__(self, num_features, num_classes, hidden_channels = 16):
        super().__init__()
        self.lin1 = torch.nn.Linear(num_features, hidden_channels)
        self.relu = torch.nn.ReLU()
        self.lin2 = torch.nn.Linear(hidden_channels, num_classes)
        self.dropout = torch.nn.Dropout(0.1)

    def forward(self, x, edge_index):
        x = self.lin1(x)
        x = self.relu(x)
        x = self.dropout(x)
        x = self.lin2(x)
        return x
```

(b) Complete the training loop below.

(2 points)

Hint: compute the loss only on the training nodes.

Hint 2: `data.x` contains the features for each node, `data.y` contains their labels.

Hint 3: `model()` takes two parameters: a tensor of node features, and a tensor of edges. See the `test_accuracy` function.

```
In [28]: device = d2l.cpu()

def accuracy(pred_y, true_y):
    correct = pred_y.argmax(dim=1) == true_y
    return int(correct.sum()) / len(true_y)

def test(model, data):
    loss_fn = torch.nn.CrossEntropyLoss()
    with torch.no_grad():
        model.eval()
        out = model(data.x, data.edge_index)
        # Compute loss and accuracy only on the 'test' nodes
        test_loss = loss_fn(out[data.test_mask], data.y[data.test_mask]).item()
        test_acc = accuracy(out[data.test_mask], data.y[data.test_mask])
```

```

        # Compute Loss and accuracy only on the 'train' nodes
        train_loss = loss_fn(out[data.train_mask], data.y[data.train_mask]).item()
        train_acc = accuracy(out[data.train_mask], data.y[data.train_mask])
        return train_loss, train_acc, test_loss, test_acc

def train(model, data, lr=0.01, weight_decay=5e-4, epochs=400, plot=True, device="c
    model = model.to(device)
    data = data.to(device)
    optimizer = torch.optim.Adam(model.parameters(), lr=lr, weight_decay=weight_dec
    loss_fn = torch.nn.CrossEntropyLoss()
    if plot:
        animator = d2l.Animator(xlabel='epoch', xlim=[1, epochs], figsize=(10, 5),
                                legend=['train loss', 'train accuracy', 'test loss'

    for epoch in range(1, epochs+1):
        model.train()
        # TODO: Compute and optimize Loss
        out = model(data.x, data.edge_index)
        loss = loss_fn(out[data.train_mask], data.y[data.train_mask])
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()
        # Compute test accuracy, and plot
        if plot and epoch % 10 == 0:
            train_loss, train_acc, test_loss, test_acc = test(model, data)
            animator.add(epoch + 1, (train_loss, train_acc, test_loss, test_acc))

    # Print final accuracy
    train_loss, train_acc, test_loss, test_acc = test(model, data)
    print(f'Train loss: {train_loss:.4f}, Train accuracy: {train_acc:.4f}')
    print(f'Test loss: {test_loss:.4f}, Test accuracy: {test_acc:.4f}')

```

(c) Now construct and train an MLP on this dataset.

(1 point)


```

In [29]: # TODO: construct and train the model
num_features = data.x.shape[1]
num_classes = len(torch.unique(data.y))
mlp_model = MLP(num_features, num_classes)
train(mlp_model, data, device='cuda')

```

Train loss: 0.0778, Train accuracy: 1.0000

Test loss: 1.2023, Test accuracy: 0.6090



(d) The MLP network does not use the citation information at all. Give a way to incorporate the edge information without using a graph neural network? (1 point)

Note that the method should still work for arbitrary citation graphs.

TODO: your answer here

6.3 A graph convolutional neural network (3 points)

Next, we will use a graph neural network based on the Graph Convolutional Network approach, which was introduced in the paper [Semi-Supervised Classification with Graph Convolutional Networks](#).

(a) Implement a graph convolutional neural network, by replacing the linear layers in the MLP with `GCNConv` layers, and train the network. (1 point)

The network should have two `GCNConv` layers. The rest of the architecture should stay as close as possible to the MLP.

```
In [30]: class GCN(torch.nn.Module):
def __init__(self, num_features, num_classes, hidden_channels = 16):
    super().__init__()
    # TODO: initialize network layers
    super().__init__()
```

```

self.conv1 = GCNConv(num_features, hidden_channels)
self.relu = torch.nn.ReLU()
self.dropout = torch.nn.Dropout(0.1)
self.conv2 = GCNConv(hidden_channels, num_classes)


def forward(self, x, edge_index):
    # TODO: compute network output
    x = self.conv1(x, edge_index)
    x = self.relu(x)
    x = self.dropout(x)
    x = self.conv2(x, edge_index)
    return x

# TODO: construct and train the model
gcn_model = GCN(num_features, num_classes)
train(gcn_model, data, device='cuda')

```

Train loss: 0.1905, Train accuracy: 1.0000

Test loss: 0.9505, Test accuracy: 0.7110



(b) Compare the results of the MLP and the GCN. Which model is better? (1 point)

TODO: your answer here

(c) Has the GCN training converged? Can you expect higher test accuracies by training longer? Explain your answer. (1 point)

TODO: your answer here

6.4 Comparing GNN layers (8 points)

Two graph layers that are interesting to compare are `SAGEConv` and `GraphConv`. Aside from one of them supporting weighted graphs, these models differ only in the accumulation function.

(a) Look at the documentation for these two layers. What is the difference in the accumulation function? (1 point)

The `GraphConv` uses as aggregation function the sum of the features of the neighbors (can be changed to max or mean), while the `SAGEConv` uses a more complex learnable aggregation function (where the default aggregator scheme is mean, but can be changed to any aggregator).

To avoid having to copy the GNN structure every time, we can make our code generic in the type of layer to use.

(b) Make a generic graph neural network, that uses layers of type `layer_type`. (1 point)

Hint: you can construct layers with `my_layer = layer_type(in_size, out_size, **layer_args)`.

```
In [31]: class GNN(torch.nn.Module):
    def __init__(self, layer_type, num_features, num_classes, hidden_channels=16, *
        super().__init__()
        # TODO: initialize network layers
        self.layer1 = layer_type(num_features, hidden_channels, **layer_args)
        self.relu = torch.nn.ReLU()
        self.dropout = torch.nn.Dropout(0.1)
        self.layer2 = layer_type(hidden_channels, num_classes, **layer_args)


    def forward(self, x, edge_index):
        # TODO: same as before
        x = self.layer1(x, edge_index)
        x = self.relu(x)
        x = self.dropout(x)
        x = self.layer2(x, edge_index)
        return x
```

(c) Train a `SAGEConv` network and a `GraphConv` network. (no points)

```
In [32]: # TODO: construct and train a GNN with SAGEConv layers
sageconv_model = GNN(SAGEConv, num_features, num_classes)
train(sageconv_model, data)
```

Train loss: 0.0503, Train accuracy: 1.0000

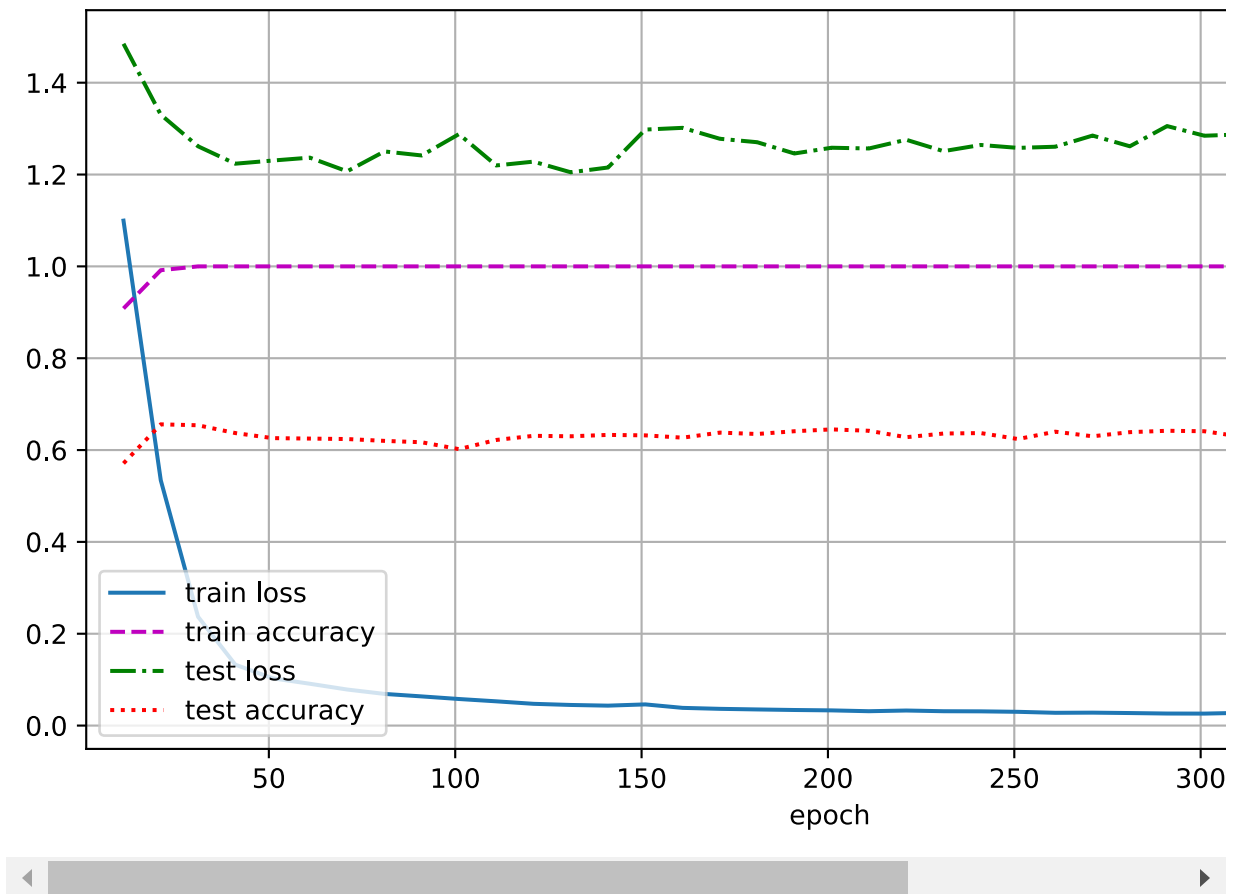
Test loss: 0.9260, Test accuracy: 0.6950



```
In [33]: # TODO: construct and train a GNN with GraphConv Layers
graphconv_model = GNN(GraphConv, num_features, num_classes)
train(graphconv_model, data)
```

Train loss: 0.0240, Train accuracy: 1.0000

Test loss: 1.3175, Test accuracy: 0.6340



(d) Compare the performance of these two models, and also compare them to the GCN. (1 point)

Hint: look at the test loss.

TODO: your answer here

(e) Can you explain the observation in the previous question by looking at the aggregation functions? Why is one of them worse than the others? (1 point)

TODO: your answer here

In fact, it is possible to use different aggregation functions, by passing `aggr=` to the network constructor.

(f) Compute the performance for GraphConv networks with 'mean', 'sum', 'min', 'max', and 'std' aggregation. (1 point)

Hint: train with `plot=False` to only show the final loss and accuracy.

Hint 2: if the performance is the same for all methods, there is most likely a bug in your GNN code.

```
In [34]: for aggr in ['mean', 'sum', 'min', 'max', 'std']:
```

```
train(GNN(GraphConv, num_features, num_classes, aggr=aggr), data, plot=False)
```

Train loss: 0.0483, Train accuracy: 1.0000
Test loss: 0.9324, Test accuracy: 0.6970
Train loss: 0.0237, Train accuracy: 1.0000
Test loss: 1.3602, Test accuracy: 0.6340
Train loss: 0.0475, Train accuracy: 1.0000
Test loss: 1.0358, Test accuracy: 0.6730
Train loss: 0.0319, Train accuracy: 1.0000
Test loss: 1.0393, Test accuracy: 0.6660
Train loss: 0.0689, Train accuracy: 1.0000
Test loss: 1.1822, Test accuracy: 0.6280

(g) Which three aggregation methods are the worst? For each one, explain why that one would not work well. (3 points)

Hint: bag-of-words features are very sparse.

TODO: your answer here

6.5 Discussion (3 points)

(a) Our training procedure gets the entire graph, including test nodes. Is it possible for the model to cheat using leaked information? (1 point)

TODO: your answer here

(b) Can the GCN and GNN networks use information from neighbors of neighbors to classify a node? Briefly explain your answer. (1 point)

TODO: your answer here

(c) Do you think the trained model will generalize to other graphs? Motivate your answer. (1 point)

TODO: your answer here

The end

Well done! Please double check the instructions at the top before you submit your results.

This assignment has 21 points.

Version f502e67 / 2023-10-04